

# Numerical Optimization: Basics

Oswin Krause, NO, 2022

KØBENHAVNS UNIVERSITET



## Quick notes on Previous Assignment

- Goal of Experimental part:

Think about how to show correctness of your code

- This is a major problem in this course: you have to convince us that your code is correct
- You have to convince yourself that your code is correct

## Quick notes on Previous Assignment

### An example solution

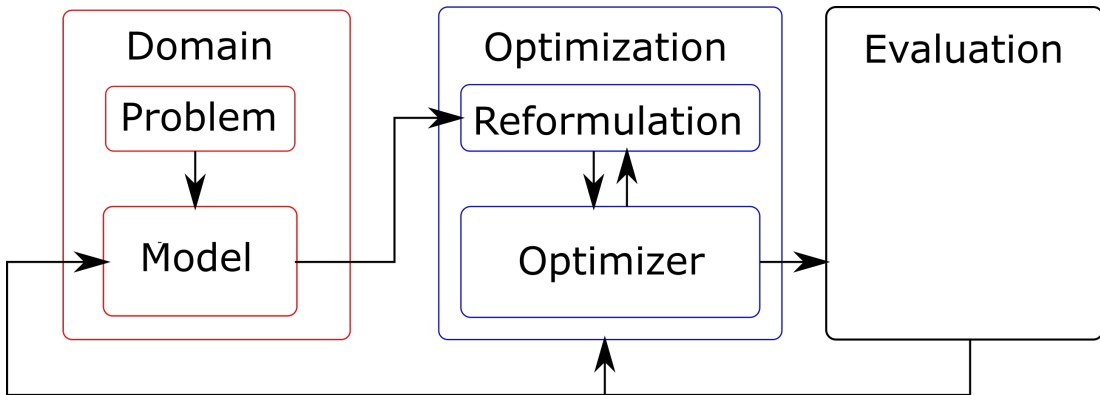
*To test my implementation, I compared my code against the gradients computed by autograd.numpy in 2D. I evaluated gradients and Hessians at the location of 100 uniformly randomly sampled points in the range  $[-2, 2]^2$ . I measured the absolute difference between the vector/matrix entries computed by my code and the corresponding entry of the ground truth computation.*

*Over all estimates, the largest deviation was  $10^{-5}$  on  $f_3$ . Compared to a value of the corresponding partial derivative of 2 at that point, this is a very small deviation.*

*I further performed the following tests...*


*I therefore conclude, that my numerical implementation is close to the result computed by autograd and that my gradients pass basic comparisons with the analytic gradients.*

## Role of Optimization



# Example: Deriving a Problem

## Example: Maximum inscribed circle problem




- Task: find the largest circle fitting into a polytope
- Modeling: Find the ball with maximum volume such that all its points are inside a convex polytope

## Example (cont): Describing the objects

Ball:

$$B(m, r) = \{x \in \mathbb{R}^n \mid \|x - m\| \leq r\}$$




## Example (cont): Describing the objects

Convex Polytope:

- $K$  Linear Inequalities

$$a_i^T x + b_i \leq 0, \quad i = 1, \dots, K$$

- Polytope  $P$  set of  $x \in \mathbb{R}^n$  fulfilling all inequalities





## Example (cont): Optimization problem (model)

$$\begin{aligned} \max_{r,m} \text{Vol}(B(m,r)) \\ \text{s.t. } x \in P, \forall x \in B(m,r) \\ \wedge r \geq 0 \end{aligned}$$

- Objective: maximize volume of the ball
- Feasible region: Only allow values for  $(m,r)$  which lead to balls contained in  $P$
- This formulation is a good model, but can't be computed
- Needs reformulation!

## Example (cont): Reformulation objective

- Reformulate  $\text{Vol}(B(m, r))$

## Example (cont): Reformulation objective

- Reformulate  $\text{Vol}(B(m, r))$
- Make use of closed form formula
  - $n = 2$ : volume=area

$$\text{Vol}(B(m, r)) = 2\pi r^2$$

- Does not depend on  $m$

## Example (cont): Reformulation objective

- Reformulate  $\text{Vol}(B(m, r))$
- Make use of closed form formula
  - $n = 2$ : volume=area

$$\text{Vol}(B(m, r)) = 2\pi r^2$$

- Does not depend on  $m$
  - Since  $r > 0$ 
    - Objective monotonic increasing in  $r$
- Optimal  $r$  does not change if we just maximize  $r$

## Example (cont): Optimization problem (intermediate)

$$\begin{aligned} \max_{r,m} \quad & r \\ \text{s.t.} \quad & x \in P, \forall x \in B(m, r) \\ & \wedge r \geq 0 \end{aligned}$$

- Need to reformulate constraint

## Example (cont): Reformulation constraint

- Goal: only allow balls that are fully contained in the polytope

## Example (cont): Reformulation constraint

- Goal: only allow balls that are fully contained in the polytope
- Idea:
  - Check that midpoint is inside
  - And its minimum distance to any point on boundaries is  $\geq r$

## Example (cont): Reformulation constraint

- Goal: only allow balls that are fully contained in the polytope
- Idea:
  - Check that midpoint is inside
  - And its minimum distance to any point on boundaries is  $\geq r$
- Both conditions are fulfilled if (assuming  $\|a_i\| = 1$ )

$$a_i^T m + b_i \leq -r, \text{ for } i = 1, \dots, K$$

- (We might get back to that later, Week 7.)



## Optimization problem (final)

$$\begin{aligned} \max_{r,m} \quad & r \\ \text{s.t.} \quad & a_i^T m + b_i \leq -r, \text{ for } i = 1, \dots, K \\ & \wedge r \geq 0 \end{aligned}$$

This is a linear program

- Objective function is linear
  - All constraints are linear
- There are specialized algorithms for this type of problem (super fast!)

# Example: Optimization Basics

## Setup in this course

- We are given an *objective function*  $f : \mathbb{R}^n \rightarrow \mathbb{R}$ .
- A *feasible region*,  $\mathcal{C} \subseteq \mathbb{R}^n$
- *Feasible regions* are represented by *constraints*
- Task: Search for a point  $x^* \in \mathcal{C}$  where  $f$  is minimal

$$x^* = \arg \min_{x \in \mathcal{C}} f(x)$$

# Numerical optimization

Goal of numerical optimization:

For a given objective function  $f$  with feasible region  $\mathcal{C}$ , find an optimization algorithm  $\mathcal{A}$  that produces a series of candidate points  $x_k$  such, that

- Often  $f(x_{k+1}) < f(x_k)$
- $(x_k)_k$  converges to some  $x' \in \mathcal{C}$
- $x'$  is a minimizer

# Numerical optimization

Goal of numerical optimization:

For a given objective function  $f$  with feasible region  $\mathcal{C}$ , find an optimization algorithm  $\mathcal{A}$  that produces a series of candidate points  $x_k$  such, that

- Often  $f(x_{k+1}) < f(x_k)$
- $(x_k)_k$  converges to some  $x' \in \mathcal{C}$
- $x'$  is a minimizer

This is very abstract. We need to know:

- How do we recognize a minimizer?
- How do we know  $(x_k)_k$  approaches a minimizer?
- How do we measure how well an algorithm does that?
- What properties do we expect of an algorithm?

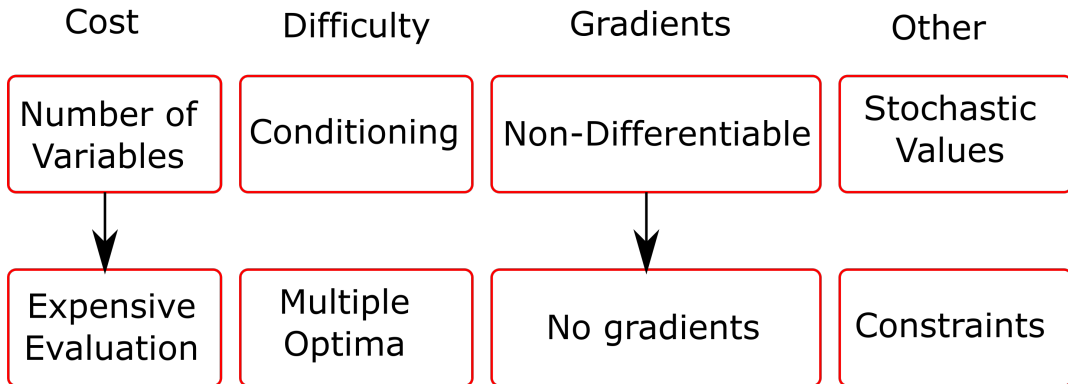
## Desirable properties of an algorithm

We want  $\mathcal{A}$  to be

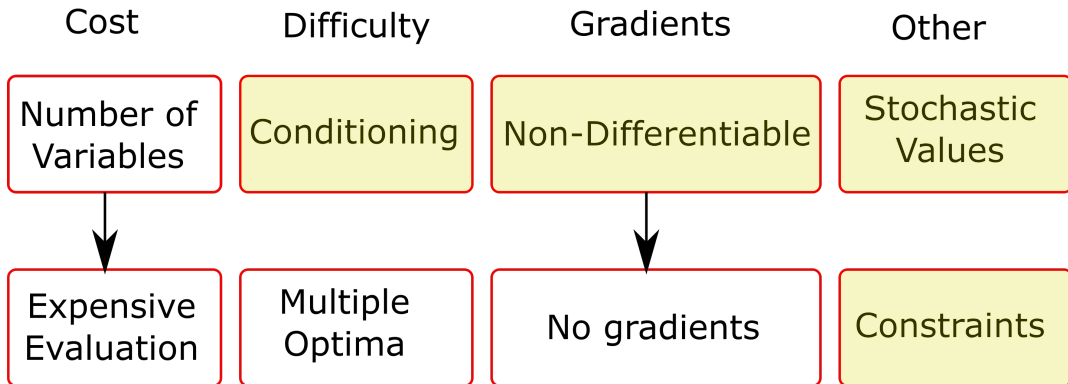
- Robust: Works on as many  $f$  as possible
- Efficient: Finds the optimum quickly with small computation cost
- Accurate: Locates the optimum with high precision
- Stable: Numeric implementation does not fail

Discuss (8 min): How can we show that an algorithm is robust, efficient and/or accurate and stable? Are they mutually exclusive?

## Difficulties when optimizing a function



## Difficulties when optimizing a function





# Unconstrained Continuous Optimization

## Setup

For now, we restrict ourselves to simple problems:

- $f : \mathbb{R}^n \rightarrow \mathbb{R}$  is a smooth function
- smooth: infinitely often continuous differentiable.
- Unconstrained (Feasible region  $\mathcal{C} = \mathbb{R}^n$ )
  - For convergence proofs of algorithms, we might need stronger conditions

# Minimisers I

## Definition (Minimisers)

1.  $x^*$  is a weak local minimiser of  $f$  if there is a neighbourhood  $\mathcal{N} \subseteq \mathbb{R}^n$  of  $x^*$  with

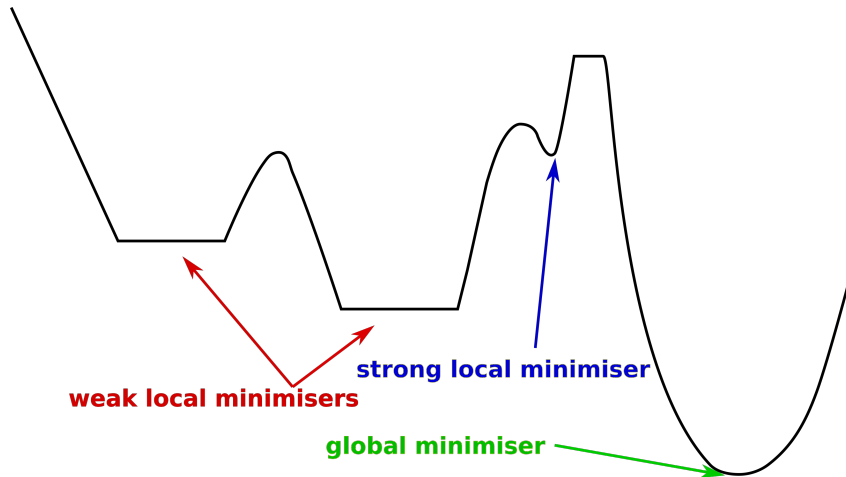
$$\forall x \in \mathcal{N}, f(x^*) \leq f(x)$$

2.  $x^*$  is a strict/strong local minimiser of  $f$  if there is a neighbourhood  $\mathcal{N} \subseteq \mathbb{R}^n$  of  $x^*$  with

$$\forall x \in \mathcal{N} \setminus \{x^*\}, f(x^*) < f(x)$$

3.  $x^*$  is an isolated local minimiser of  $f$  if there is a neighbourhood  $\mathcal{N} \subseteq \mathbb{R}^n$  of  $x^*$  such that  $x^*$  is the unique minimiser of  $f$  in  $\mathcal{N}$ .

## Minimisers II



# Conditions for Minimizers

## Theorem (Necessary conditions)

- *First order necessary condition. If  $x^*$  is a local minimiser of  $f$  over a neighbourhood  $\mathcal{N}$ , then  $\nabla f(x^*) = 0$ , i.e,  $x^*$  is a critical point of  $f$ .*
- *Second order necessary condition. At a critical point,  $\nabla^2 f(x^*)$  must be positive semidefinite.*

## Theorem (Second order sufficient conditions)

*Assume that  $x^*$  is a critical point of  $f$ . Then if  $\nabla^2 f(x^*)$  is positive definite,  $x^*$  is a strict local minimiser*

## Proof: First Order Necessary Conditions

Proof idea: Show that we can find a better point in any neighbourhood around  $x^*$ , when  $\nabla f(x^*) \neq 0$

## Proof: First Order Necessary Conditions

Proof idea: Show that we can find a better point in any neighbourhood around  $x^*$ , when  $\nabla f(x^*) \neq 0$

- Assume  $\nabla f(x^*) \neq 0$

## Proof: First Order Necessary Conditions

Proof idea: Show that we can find a better point in any neighbourhood around  $x^*$ , when  $\nabla f(x^*) \neq 0$

- Assume  $\nabla f(x^*) \neq 0$
- Let  $p = -\nabla f(x^*)$ . We have

$$p^T \nabla f(x^*) = -\nabla f(x^*)^T \nabla f(x^*) = -\|\nabla f(x^*)\|^2 < 0$$



## Proof: First Order Necessary Conditions

Proof idea: Show that we can find a better point in any neighbourhood around  $x^*$ , when  $\nabla f(x^*) \neq 0$

- Assume  $\nabla f(x^*) \neq 0$
- Let  $p = -\nabla f(x^*)$ . We have

$$p^T \nabla f(x^*) = -\nabla f(x^*)^T \nabla f(x^*) = -\|\nabla f(x^*)\|^2 < 0$$

- Define  $g(t) = f(x^* + tp) - f(x^*)$  with
  - $g'(t) = \nabla f(x^* + tp)^T p$
  - $g'(0) = -\|\nabla f(x^*)\|^2 < 0$

## Proof: First Order Necessary Conditions

Proof idea: Show that we can find a better point in any neighbourhood around  $x^*$ , when  $\nabla f(x^*) \neq 0$

- Assume  $\nabla f(x^*) \neq 0$
- Let  $p = -\nabla f(x^*)$ . We have

$$p^T \nabla f(x^*) = -\nabla f(x^*)^T \nabla f(x^*) = -\|\nabla f(x^*)\|^2 < 0$$

- Define  $g(t) = f(x^* + tp) - f(x^*)$  with
  - $g'(t) = \nabla f(x^* + tp)^T p$
  - $g'(0) = -\|\nabla f(x^*)\|^2 < 0$
- Statement true, if  $\exists T > 0$  such, that for all  $0 < t \leq T$  holds  $g(t) < 0$

## Proof: First Order Necessary Condition

First: Use Taylor expansion on  $g$  around 0 to introduce the derivative

- We have  $g(0) = 0$  and  $g'(0) < 0$
- Remember:  $f$  (and thus  $g$  are smooth)

## Proof: First Order Necessary Condition

First: Use Taylor expansion on  $g$  around 0 to introduce the derivative

- We have  $g(0) = 0$  and  $g'(0) < 0$
- Remember:  $f$  (and thus  $g$ ) are smooth)

→ We can create first-order Taylor expansion of  $g$  around 0

$$g(0+t) = \underbrace{g(0)}_0 + t \underbrace{g'(0)}_{<0} + \underbrace{\frac{t^2}{2} g''(c)}_{\text{Error Term}}, \quad c \in [0, t]$$

- Remember:  $c$  depends on the chosen  $t$ .

## Proof: First Order Necessary Condition

First: Use Taylor expansion on  $g$  around 0 to introduce the derivative

- We have  $g(0) = 0$  and  $g'(0) < 0$
- Remember:  $f$  (and thus  $g$ ) are smooth)

→ We can create first-order Taylor expansion of  $g$  around 0

$$g(0+t) = \underbrace{g(0)}_0 + t \underbrace{g'(0)}_{<0} + \underbrace{\frac{t^2}{2} g''(c)}_{\text{Error Term}}, \quad c \in [0, t]$$

- Remember:  $c$  depends on the chosen  $t$ .
- Next step: pick  $t$  small enough that the error term is smaller than  $tg'(0)$

## Proof: First Order Necessary Condition

Final step: construct points with lower function value.

- Let  $M = \max_{c \in [0,1]} |g''(c)|$ .

## Proof: First Order Necessary Condition

Final step: construct points with lower function value.

- Let  $M = \max_{c \in [0,1]} |g''(c)|$ .
- Pick  $T = \min\{-\frac{g'(0)}{M}, 1\}$
- We have for  $T < 1$

$$Tg'(0) + \frac{T^2}{2} \underbrace{\overbrace{g''(c)}^{c \in [0,1]}}_{\in [-M, M]} = -\frac{(g'(0))^2}{M} + \frac{1}{2} \frac{(g'(0))^2}{M} \underbrace{\frac{g''(c)}{M}}_{\in [-1,1]} < 0$$

## Proof: First Order Necessary Condition

Final step: construct points with lower function value.

- Let  $M = \max_{c \in [0,1]} |g''(c)|$ .
- Pick  $T = \min\{-\frac{g'(0)}{M}, 1\}$
- We have for  $T < 1$

$$Tg'(0) + \frac{T^2}{2} \underbrace{\overbrace{g''(c)}^{c \in [0,1]}}_{\in [-M, M]} = -\frac{(g'(0))^2}{M} + \frac{1}{2} \frac{(g'(0))^2}{M} \underbrace{\frac{g''(c)}{M}}_{\in [-1,1]} < 0$$

- For  $T = 1$ , the results holds, since  $|g''(c)| \leq |g'(0)|$
- As the error terms shrinks faster with  $t$  than the linear term, this holds  $\forall t \in (0, T]$ .

$\Rightarrow$  There is no neighbourhood around  $x^*$  where  $f(x^*)$  is minimal

$\Rightarrow x^*$  cannot be local minimum.



## Special case: Convex functions

### Definition (Convex Functions)

Let  $f : \mathbb{R}^n \rightarrow \mathbb{R}$  be a function. We call  $f$  *convex* if  $\forall x, p \in \mathbb{R}^n$  and  $\lambda \in (0, 1)$  holds

$$f(x + \lambda p) \leq f(x) + \lambda(f(x + p) - f(x)) .$$

Further, if the stricter condition

$$f(x + \lambda p) < f(x) + \lambda(f(x + p) - f(x)) .$$

holds, then we call  $f$  *strictly convex*

Note: convexity is often written by substitution  $y = x + p$ . This formulation above is more useful to us.

# Characterization of differentiable convex functions

## Lemma

*Let  $\mathcal{C} \subseteq \mathbb{R}^n$  and  $f : \mathcal{C} \rightarrow \mathbb{R}$  be a twice continuous differentiable function.*

*$f$  is convex iff*

*$\nabla^2 f(x)$  is positive semi-definite*

*(for strict convexity,  $\nabla^2 f(x)$  is PD)*

## Why convexity matters

- One can show:  $f$  convex  $\Rightarrow f$  has at most one local minimum.

## Why convexity matters

- One can show:  $f$  convex  $\Rightarrow f$  has at most one local minimum.
- For all smooth  $f$ 
  - Second order necessary condition of minima:  $\nabla^2 f$  is PD at a minimum

## Why convexity matters

- One can show:  $f$  convex  $\Rightarrow f$  has at most one local minimum.
- For all smooth  $f$ 
  - Second order necessary condition of minima:  $\nabla^2 f$  is PD at a minimum
  - Continuity:  $f$  is convex in a neighbourhood close to a local minimum
  - Algorithms that work well on convex functions, work well close to the optimum
  - We just have to ensure to get there.

# Basics of Algorithm Design and Evaluation

## Basic Idea of optimization algorithm

- Target function  $f$  complex, hard to compute
- We can't find minima directly

## Basic Idea of optimization algorithm

- Target function  $f$  complex, hard to compute
- We can't find minima directly
- Idea of basic optimization loop
    1. Approximate  $f(x + p)$  around  $x$  by simpler function
      - $m : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $m(p) \approx f(x + p)$
      - Ensure that  $m$  can be analyzed easily



# Basic Idea of optimization algorithm

- Target function  $f$  complex, hard to compute
- We can't find minima directly
- Idea of basic optimization loop
    1. Approximate  $f(x + p)$  around  $x$  by simpler function
      - $m : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $m(p) \approx f(x + p)$
      - Ensure that  $m$  can be analyzed easily
    2. Find  $p$  with  $m(p) < m(0) \cong f(x)$

# Basic Idea of optimization algorithm

- Target function  $f$  complex, hard to compute

→ We can't find minima directly

- Idea of basic optimization loop
  1. Approximate  $f(x + p)$  around  $x$  by simpler function
    - $m : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $m(p) \approx f(x + p)$
    - Ensure that  $m$  can be analyzed easily
  2. Find  $p$  with  $m(p) < m(0) \cong f(x)$
  3. Evaluate  $f(x + p)$  and correct  $p$  (or  $m$ ) to ensure  $f(x + p) < f(x)$

## Basic Idea of optimization algorithm

- Target function  $f$  complex, hard to compute

→ We can't find minima directly

- Idea of basic optimization loop

1. Approximate  $f(x + p)$  around  $x$  by simpler function

- $m : \mathbb{R}^n \rightarrow \mathbb{R}$ ,  $m(p) \approx f(x + p)$
- Ensure that  $m$  can be analyzed easily

2. Find  $p$  with  $m(p) < m(0) \cong f(x)$

3. Evaluate  $f(x + p)$  and correct  $p$  (or  $m$ ) to ensure  $f(x + p) < f(x)$

4. Set  $x \rightarrow x + p$  and go to 1.

## Example: Line-Search based Gradient Descent

1. Set  $m(p) = f(x) + p^T \nabla f(x)$  (first order Taylor)
2. Pick  $p = -\nabla f(x)$
3. Find  $\alpha$  such, that  $f(x + \alpha p) < f(x)$
4. Set  $x \rightarrow x + \alpha p$  and go to 1.

## Example: Trust-Region Newton

1. Set  $m(p) = f(x) + p^T \nabla f(x) + \frac{1}{2} p^T \nabla^2 f(x) p$  (second order Taylor)
2.  $p = \min_p m(p)$  so, that  $\|p\| \leq \Delta$
3. If  $f(x + p) > f(x)$ 
  - 3.1 Reduce  $\Delta$
  - 3.2 Go to 2.
4. Set  $x \rightarrow x + p$  and go to 1.

## When to stop?

- Let  $(x_k)_k$  be the sequence of iterates by the algorithm
- In most cases  $f(x^*) \neq 0$  and  $x^*$  unknown
- How do we know, when to stop? Discuss (5min)

## When to stop?

- Let  $(x_k)_k$  be the sequence of iterates by the algorithm
- In most cases  $f(x^*) \neq 0$  and  $x^*$  unknown
- How do we know, when to stop? Discuss (5min)
- Useful criteria
  - Maximum number of iterations
  - Length of steps(?)
  - Improvement of function value(?)
  - Norm of the gradient

## When to stop?

- Let  $(x_k)_k$  be the sequence of iterates by the algorithm
- In most cases  $f(x^*) \neq 0$  and  $x^*$  unknown
- How do we know, when to stop? Discuss (5min)
- Useful criteria
  - Maximum number of iterations
  - Length of steps(?)
  - Improvement of function value(?)
  - Norm of the gradient
- Norm of the gradient?
  - We know  $(\nabla f(x_k))_k \rightarrow 0$ , when  $(x_k)_k \rightarrow x^*$   
→ use  $\|\nabla f(x_k)\| < \epsilon$  as criterion



## Measuring convergence Rate

- As algorithm designers we want to know how good algorithms are.
- We can measure their performance on Benchmark functions.
- Discuss: What are good metrics to compare algorithms? (5min)

## Measuring convergence Rate

- As algorithm designers we want to know how good algorithms are.
- We can measure their performance on Benchmark functions.
- Discuss: What are good metrics to compare algorithms? (5min)
- Often used:
  - Number of steps until termination
  - Function value after  $N$  steps
- Problem: only tells us about performance at one point
  - Algorithms might slow down over time
  - Or speed up

## Measuring convergence Rate

- As algorithm designers we want to know how good algorithms are.
- We can measure their performance on Benchmark functions.
- Discuss: What are good metrics to compare algorithms? (5min)
- Often used:
  - Number of steps until termination
  - Function value after  $N$  steps
- Problem: only tells us about performance at one point
  - Algorithms might slow down over time
  - Or speed up
- Can we somehow measure the convergence rate of the sequences:
  - $(\|x_k - x^*\|)_k$
  - $(f(x_k) - f(x^*))_k$
  - $(\|\nabla f(x_k)\|)_k$

## Experimentally verifying Linear Convergence

Remember linear Convergence (example with  $\|x_k - x^*\|$ ):

$$\frac{\|x_k - x^*\|}{\|x_{k-1} - x^*\|} \leq M$$

## Experimentally verifying Linear Convergence

Remember linear Convergence (example with  $\|x_k - x^*\|$ ):

$$\frac{\|x_k - x^*\|}{\|x_{k-1} - x^*\|} \leq M$$

Rewrite:

$$\begin{aligned}\|x_k - x^*\| &\leq M\|x_{k-1} - x^*\| \\ &\leq M^2\|x_{k-2} - x^*\| \\ &\leq \dots \\ &\leq M^k\|x_0 - x^*\|\end{aligned}$$

## Experimentally verifying Linear Convergence

Remember linear Convergence (example with  $\|x_k - x^*\|$ ):


$$\frac{\|x_k - x^*\|}{\|x_{k-1} - x^*\|} \leq M$$

Rewrite:


$$\begin{aligned}\|x_k - x^*\| &\leq M\|x_{k-1} - x^*\| \\ &\leq M^2\|x_{k-2} - x^*\| \\ &\leq \dots \\ &\leq M^k\|x_0 - x^*\|\end{aligned}$$

$\log\|x_k - x^*\| \leq k \log M + \log\|x_0 - x^*\|$  is approximately a linear function, when plotted on semi-log plot

## Linear convergence rate: line in semi-log plot



## Other convergence rates: visible





## Quadratic convergence: obvious

